
Cloud Files Introduction Documentation

Release 1.1.0

Ranjit Nayak

August 12, 2009

CONTENTS

1	Introduction	1
2	Ideal uses of Cloud Files	3
3	Why use Cloud Files	5
4	Cloud Files Features	7
4.1	Simplicity	7
4.2	Affordability	7
4.3	Flexibility	7
4.4	Superior performance	8
4.5	World Class support	8
4.6	Data Redundancy	8
4.7	Security	8
5	Tools and Applications for Cloud Files	11
5.1	Control Panel	11
5.2	FireUploader	12
5.3	Cyberduck	12
6	Developing applications for Cloud Files	15
6.1	ReST API	15
6.2	Language specific API	15
6.3	Caveats	15
7	Additional Resources	17
8	Trademarks	19

INTRODUCTION

The Rackspace Cloud Files™ is an affordable, redundant, scalable, and dynamic storage service offering. The core storage system is designed to provide a safe, secure, automatically re-sizing and network accessible way to store data. You can store an unlimited quantity of files and each one can be as large as 5 gigabytes. Users can store as much as they want and pay only for storage space they actually use.

Additionally, Cloud Files provides a simple yet powerful way to publish and distribute content behind the industry leading Limelight Networks' Content Distribution Network. Cloud Files users get access to this network automatically without having to worry about contracts, additional costs, or technical hurdles.

Cloud Files allows users to store/retrieve files and CDN-enable content via a simple Web Service (ReST: Representational State Transfer) interface. There are also language-specific API's that utilize the ReST API but make it much easier for developers to integrate into their applications.

For more details on the Cloud Files service please refer to http://www.rackspacecloud.com/cloud_hosting_products/files

IDEAL USES OF CLOUD FILES

There are a number of uses for the Cloud Files service. Cloud Files is an excellent storage solution for a number of scenarios and is well suited for a number of applications such as:

- Backing up or archiving data
- Serving images/videos (streaming data to the user's browser)
- Serving content with a world-class CDN (Limelight Networks)
- Storing secondary/tertiary, static web-accessible data
- Developing new applications with data storage integration
- Storing data when predicting storage capacity is difficult
- Storing data for applications affordably

WHY USE CLOUD FILES

There are a number of reasons to choose the Cloud Files service over similar services available in the market.

- Simplicity of security mechanisms and CDN functionality
- Affordability of the service with a price of just 15cents / gigabyte of storage
- Flexibility in terms paying as you go
- Several developer programming interfaces
- Great performance
- World class support

CLOUD FILES FEATURES

4.1 Simplicity

Cloud Files is simple to use for developers and non-developers alike. Users can get started in as little as five minutes. Users do not have to know how to code to use Cloud Files and CDN. Users can, within minutes, sign up for Cloud Files, create a Container, upload a file and publish that Container's content through the CDN (Refer to "Cloud Files with CDN QuickStart guide". The Cloud Files GUI is easy to navigate and use. Rackspace browser based control panel let users easily upload files and distribute on a CDN without writing any code.

All the content can be backed by a CDN without complex negotiations and details of updating content for optimizing delivery. Given the complexity of the CDN, users may have to make a series of choices such as number of servers to use etc. before obtaining the service. After making the choices users must ensure the usage bills are accurate. All these issues are handled by Cloud Files. CDNs have, in the past, been the prerogative of companies with more money but Cloud Files has changed that.

Cloud Files is a user and developer friendly service. Security mechanisms described in more detail in the Security section are simple to use. Third party tools which further simplify use of the storage service are available. Developers can use a language specific application programming interfaces to develop utilities or applications. The API's are easy to use and are documented with examples to get started quickly.

4.2 Affordability

Pricing for Cloud Files with CDN starts at 15 cents per gigabyte of storage and 8 cents per gigabyte inbound and 22 cents per gigabyte of bandwidth outbound to any edge location around the globe. There are no per requests fees for CDN. For the latest pricing information please refer to the Cloud Files web site.

4.3 Flexibility

Cloud Files is a dynamically expanding storage system very flexible and is built on the pay for use principle. Customers can use as much or little storage as necessary, while paying only for storage space used. There is no limit on total space use and individual files can be as large as 5GB. The GUI control panel allows users to check storage space used and bandwidth utilized. There are no upfront fees or contracts and end users pay only for storage space used and outgoing bandwidth that they use.

System administrators can use this for simple manual updates as well. Users do not have to know how to code to use Cloud Files and CDN. No API is required to share files. Non-developers have a simple web-interface, that can be used to quickly and easily upload data and enable CDN access. There are multiple third party tools (refer to Tools and Applications section) which make it even more flexible for users of specific environments such as the Mac OS.

Developers can use the ReST web service and language-specific API's in PHP, Python, Java, Ruby, and C#/.NET. The API's provide full support for managing content in Cloud Files and publishing content over the CDN. The API allow developers to work in the language they feel most comfortable

4.4 Superior performance

Rackspace has built networks with superior performance for years. The Limelight Networks CDN capability improves the performance further for distribution of digital content. Limelight's scale has no individual limits. Maximum transfer speed and requests per second are only limited by the overall, aggregate egress capacity of Limelight's network which, at this time is 2.2Tbps. In short, Limelight can bring data closer to end users and serve popular content faster. To test the performance, 8KB files were placed in Cloud Files. Then Gomez was used to measure end to end response times (including DNS resolution, etc.) to the content via the CDN URLs, every five minutes, over the course of one week, from various nodes throughout the world. The average response time in the US was 57ms for Cloud Files. The average response time from global locations was 107ms for Cloud Files. For more details and comparisons with similar services please refer to blog entries on the Rackspace Cloud blog. For more information on why CDNs improve performance please refer to the document titled "Content Delivery Networks with Cloud Files".

4.5 World Class support

With Cloud Files world-class free technical support is only a click away. Live support, with real people is available 24/7. Fanatical Support is built into the price. Users can get peace of mind knowing that technical support is just a phone call or online chat away – at any time of the day.

4.6 Data Redundancy

The Cloud Files storage system is designed to be highly available and fault tolerant. Cloud Files achieves client data redundancy by replicating three full copies on different storage nodes. Storage nodes are grouped in logical Zones within Rackspace datacenters. Zones are connected to redundant Internet backbone providers and reside on redundant power supplies and generators. The system has been engineered in such a way as to continue to be fully functional even in the event of a major service disruption within a Datacenter. Content on the CDN provides an additional layer of data redundancy.

4.7 Security

Cloud Files uses a number of different measures to ensure that your data is kept safe. First and foremost, all traffic between clients and the Cloud Files system uses SSL to establish a secure, encrypted communication channel. This ensures that any data (usernames, passwords, and content) cannot be intercepted and read by a third-party. Users authenticate with a valid username and API Access Key and are granted a session authentication token. These authentication tokens are used to validate all operations against Cloud Files. There is no way to terminate a valid session by the user, but the session tokens will automatically expire after 24 hours, forcing clients to resend their credentials.

The API Access Key is only available from within the Rackspace Cloud's web control panel at <http://www.rackspacecloud.com>. Users must enter their valid username and password to gain access to view the API Access Key or to generate a new key.

It is important to note that Cloud Files does not apply any transformations to data before storing it. This means that Cloud Files *will not* store data in encrypted form unless the client encrypts it prior to transmission. This allows users to select the type and level of encryption best suited for their application.

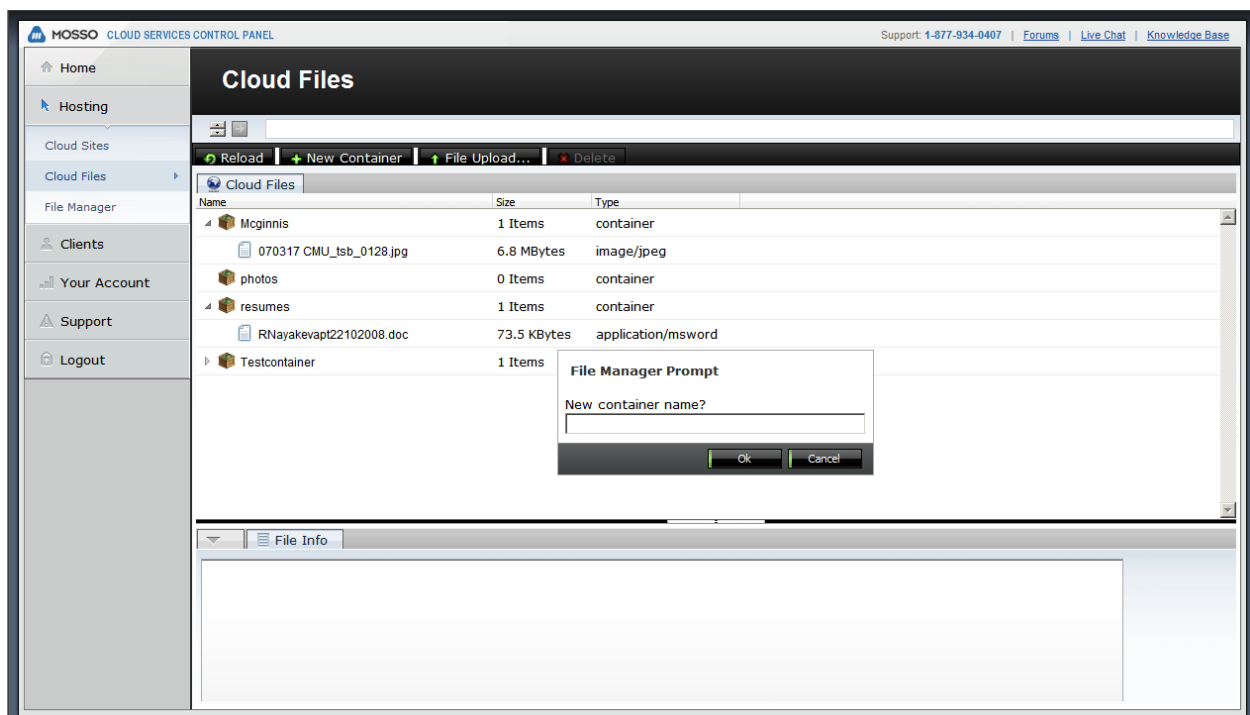
When deleting storage Objects from the Cloud Files system, all copies of data are permanently removed within a matter of minutes. Furthermore, the physical blocks making up the customer's data is zeroed over before that space is re-used for other customer data. In other words, after a delete request, the data will be unrecoverable.

TOOLS AND APPLICATIONS FOR CLOUD FILES

In addition to the ReST and language API's, there are several other ways for users to interact with Cloud Files.

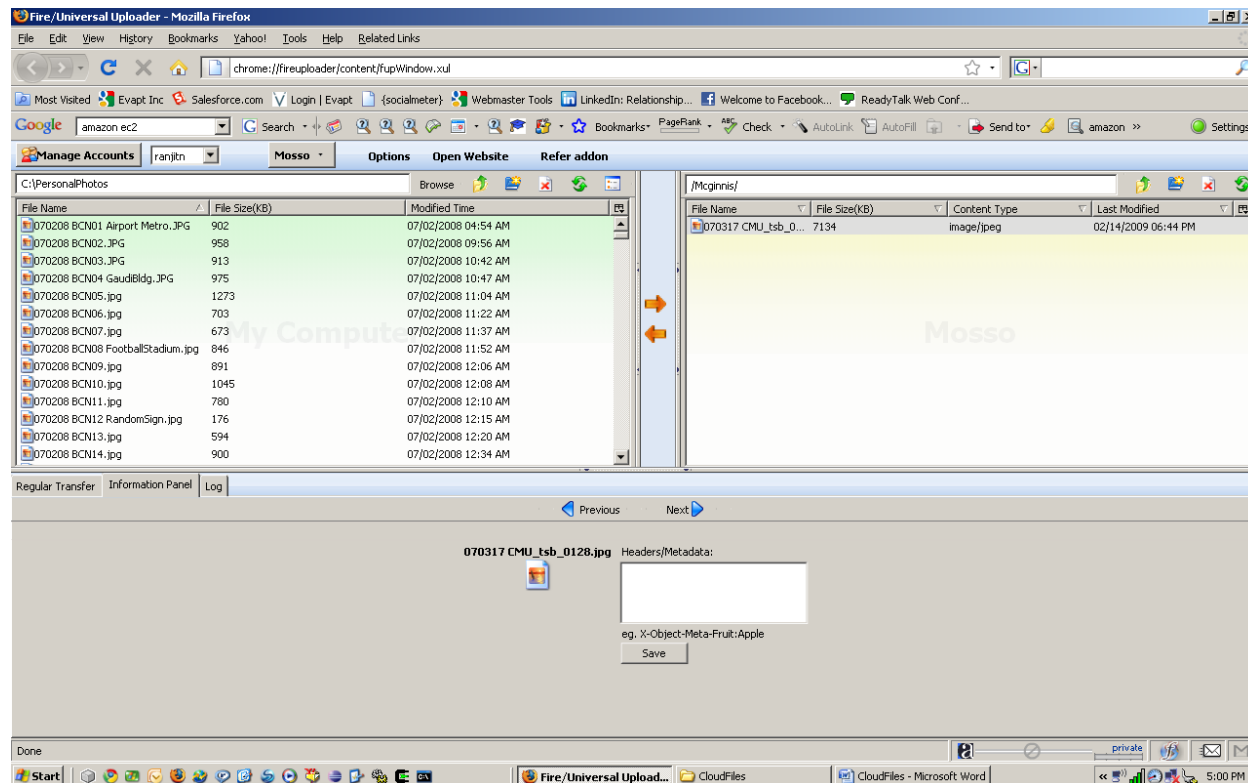
5.1 Control Panel

The control panel on The Rackspace Cloud's site is one way to access Cloud Files. Files can be uploaded and downloaded through the control panel. This is a tool to understand at a high level how Cloud Files works. It gives system administrators a view of the various containers and objects stored on Cloud Files. It is great GUI tool to update individual files and metadata on the files.



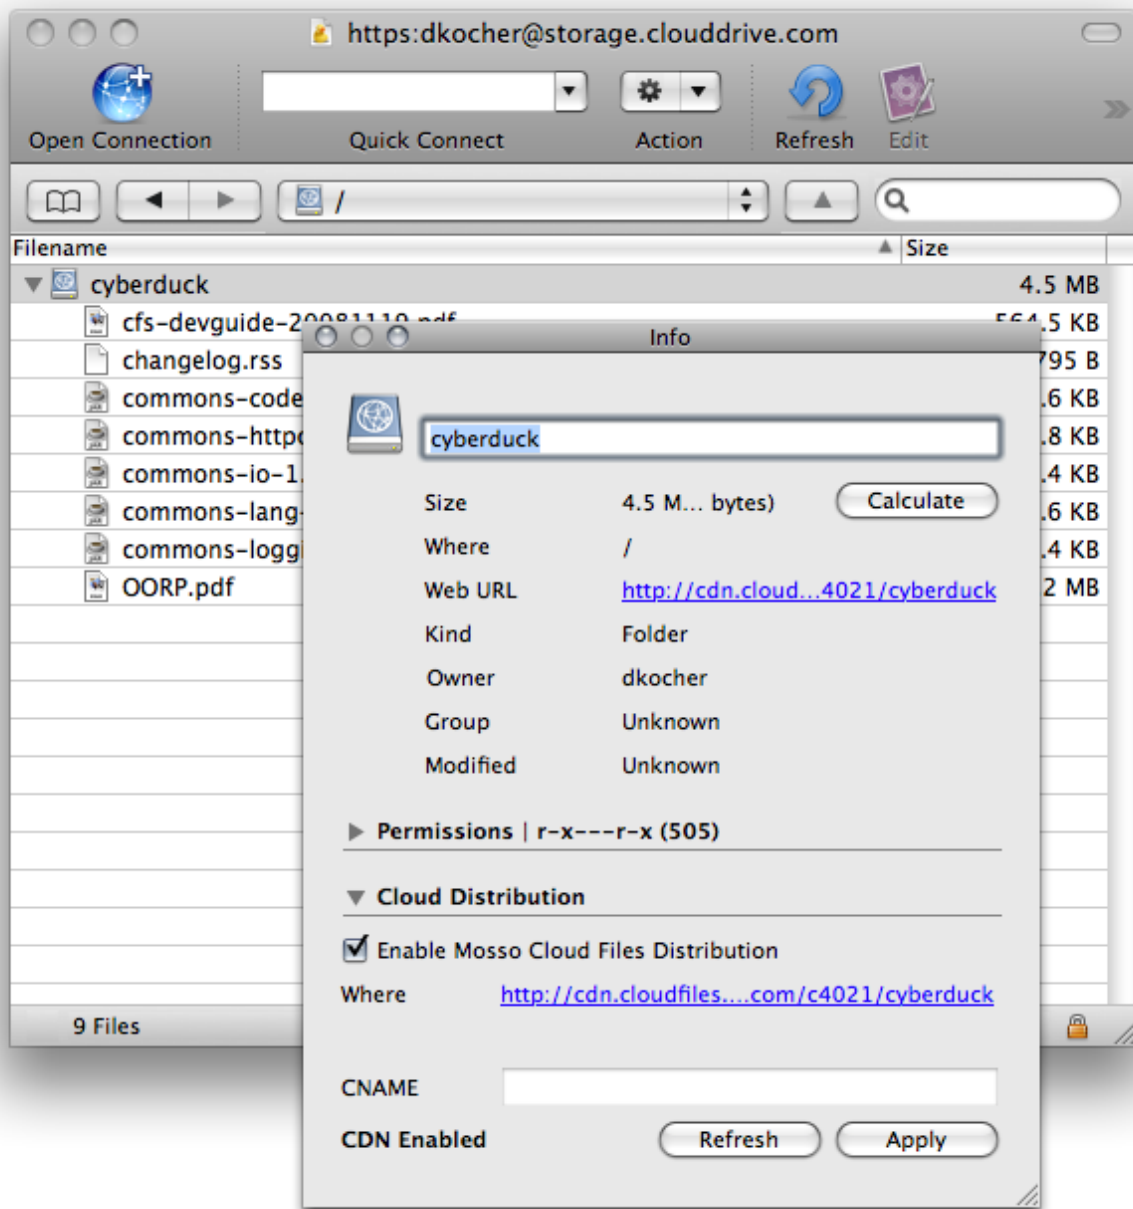
5.2 FireUploader

FireUploader is a tool from Suchi Software and it provides an intuitive interface to move to and manage files in Cloud Files. It works as a plug-in to the Mozilla Firefox browser and provides a GUI to manage the movement as well as meta data in Cloud Files as shown in Figure 2. While setting up the tool you must supply your Rackspace Cloud username and when prompted for a “password”, enter your API Access Key. This tool provides creative way of working with available Cloud Files features to allow users to simulate a hierarchical file system structure. For more information, please see the [Fireuploader web site](#).



5.3 Cyberduck

Cyberduck is an open source Cloud Files browser licensed under the GPL that runs on Macs. It features an easy to use interface with easily accessible bookmarks. The browser outline view allows users to browse large folder structures efficiently and preview files with Quick Look. Cyberduck also provides a seamless integration with many external editors making it easy to change content quickly. It provides a way to upload and configure Cloud Files to distribute content over the CDN. Cyberduck also supports many OS X core system technologies such as Spotlight, Bonjour, Keychain and a large number of language translations. Cyberduck 3.x requires Mac OS X 10.4 or later. For more information, please see the [Cyberduck web site](#).



DEVELOPING APPLICATIONS FOR CLOUD FILES

There are several programming interfaces for using Cloud Files that will allow developers to integrate the storage solution into new applications, or provide automated ways of accessing the system. The best way to get started is to follow the steps enumerated in the QuickStart guide. For more details on developing applications please refer to “Cloud Files Developer Guide”.

6.1 ReST API

Cloud Files allows users to store/retrieve files via a simple Web Service (ReST: Representational State Transfer) interface. New Cloud Files storage applications such as the tools described in the previous sections can be designed and developed by using the ReST application programming interface. The complete API documentation can be found on the Rackspace Rackspace Cloud web site.

6.2 Language specific API

In addition to the ReST API, there are language specific API documents and assets for C#/.NET, Python, PHP, Java and Ruby. For more information please refer to Developer resources on the Cloud Files web site.

6.3 Caveats

There are many great use-cases and advantages to a storage system like Cloud Files. However, there are limitations depending on your needs and some use-cases where Cloud Files is not an ideal storage solution. Keep the following points in mind as you further explore Cloud Files:

- Native support within your Operating System. It is not possible to “mount” or “map” the Cloud Files storage system as a virtual hard-disk on your computer or server.
- Disk mirroring or backup solutions that require byte/block level differences. There are no concepts of “appending data” or “file locking” operations within Cloud Files.
- Data can be organized into storage compartments called “Containers”, but Containers cannot be nested. Since there is only one top-level of organization, you will not be able to upload a nested directory/folder structures into Cloud Files unless a transformation is performed to flatten the structure.
- There is no built-in concept of permissions or access controls the Cloud Files system that you are probably used to in a traditional file system.

- There are no transcoding capabilities to read a file in one format and upload the same into one or more formats into Cloud Files.

ADDITIONAL RESOURCES

The official support channels (phone, chat, email, forums, and knowledge base articles) for Cloud Files are available at the Rackspace Cloud's web site: <http://www.rackspacecloud.com>.

There is a temporary system status page available at <http://status.cloudfiles.rackspacecloud.com> that can be reviewed if you believe the system is not functioning to your expectations. This page is updated to reflect up-to-date information about the system's current health and status. Over time, Cloud Files users will be directed to a more comprehensive status page that reflects the system health of all major cloud products.

Interested users can also follow updates/announcements via twitter at <http://www.twitter.com/cloudfiles>.

IRC savvy developers are also encouraged to join some of the members of the Cloud Files team at irc.freenode.net on the #cloudfiles channel. This is not an official Cloud Files support channel but should rather be viewed as a community meeting place to share/discuss Cloud Files.

TRADEMARKS

The following terms are trademarks of Rackspace Hosting Inc in the United States:

- Cloud Files
- Cloud Servers
- Cloud Sites

Java™ and all Java-based trademarks are trademarks of Sun Microsystems, Inc. in the United States, other countries, or both.

Linux® is a registered trademark of Linus Torvalds in the United States, other countries, or both.

UNIX® is a registered trademark of The Open Group in the United States and other countries.

Other company, product, and service names may be trademarks or service marks of others.